

INTER-PROCESS COMMUNICATION ON STEROIDS

ATLAS SOFTWARE & COMPUTING WORKSHOP

OCTOBER 2012

Roberto A. Vitillo (LBNL)

CONTEXT

- AthenaMP communications
 - ▶ reader process sends events to workers
- Coprocessor communications
 - ▶ Athena[MP] jobs interacting with a GPU server process
- Available IPC mechanisms
 - ▶ shared memory with explicit synchronization
 - ▶ message passing with implicit synchronization

MESSAGE PASSING MODEL



- One of the most successful models for providing concurrency
 - data and synchronization in a single unit
- Actor Model
 - processes have an identity
 - communicate by sending messages to mailing addresses
 - Erlang, Scala
- Process calculi
 - processes are anonymous
 - communicate by sending messages through named channels
 - Go Programming Language

PATTERNS

- Producer & Consumer
 - ▶ producer pushes messages
 - ▶ consumer pulls messages
- Client & Server
 - ▶ client makes a request
 - ▶ server replies to a request

CHANNELS

- Properties of a channels:
 - name
 - context (thread, local-process, distributed-process)
 - asynchronous(k)
 - topology



SOCKETS


- Each end of a channel is attached to a Socket
- Different patterns have different Sockets,
 - e.g. ProducerSocket, ConsumerSocket
- A Socket allows to:
 - send() buffers of data to its peers (buffer-blocking)
 - receive() buffers of data from its peers (blocking)



SOCKETS

```
Channel channel("service", ONE_TO_ONE)
ISocket *socket = factory->createClientSocket(channel);


socket->send("ping", 5);
socket->receive(&buffer);
```



```
Channel channel("service", ONE_TO_ONE);
ISocket *socket = factory->createServerSocket(channel);

while(true){
    socket->receive(&buffer);
    socket->send("pong");
}
```


DATA TRANSFER TECHNIQUES

-  Zero-Copy
 - ▶ page table entries are transferred from the source to the destination process
 - ▶ requires the buffers to have the same alignment
- Double copy in shared memory
 - ▶ double buffering allows the sender and receiver to transfer data in parallel
- Delegate the copy to a NIC
 - ▶ avoids cache pollution

DATA TRANSFER TECHNIQUES


- Single Copy with ptrace()
 - ▶ debugging facility, sender processes is stopped
 - ▶ receiver process accesses the sender's process address space and performs a copy
- Single Copy through a Kernel facility
 - ▶ vmsplice (Kernel 2.6.17)
 - ▶ KNEM (Kernel 2.6.15, requires module)
 - ▶ Cross Memory Attach (Kernel 3.2)

IMPLEMENTATION

- The API is currently implemented with ZeroMQ
 - ▶ provides a default fall back implementation
 - ▶ lock-free queues for threads
 - ▶ AF_UNIX sockets for local processes
 - ▶ TCP sockets for distributed processes
- The implementation switches according to the channel configuration
 - ▶ many-to-many channel triggers the creation of a broker process that handles all communications
 - ▶ one-to-one, producer-consumer uses a UNIX pipe with vmsplice()

BENCHMARK: LATENCY

- Test: 1 Byte Ping Pong
- MPI is using shared memory
 - double copy
 - avoids expensive system call which dominates for small messages
- The pipe implementations are dominated by the system call overhead



Ivy Bridge, ZMQ 2.2, OpenMPI 1.6.2 without KNEM


BENCHMARK: BANDWIDTH

- Benchmark: 1 Megabyte Ping Pong
- Pipe performs 2 copies (U-K-U)
- Spliced pipe performs a single copy (K-U)
- OpenMPI is using shared memory as buffer and performs 2 copies
 - copies are performed in parallel on both ends
 - consumes more CPU
 - pollutes the caches
- ZMQ uses AF_UNIX sockets, i.e. two copies are performed (U-K-U)
 - all sorts of buffers in between
 - optimized for distributed environments



BENCHMARK: BANDWIDTH

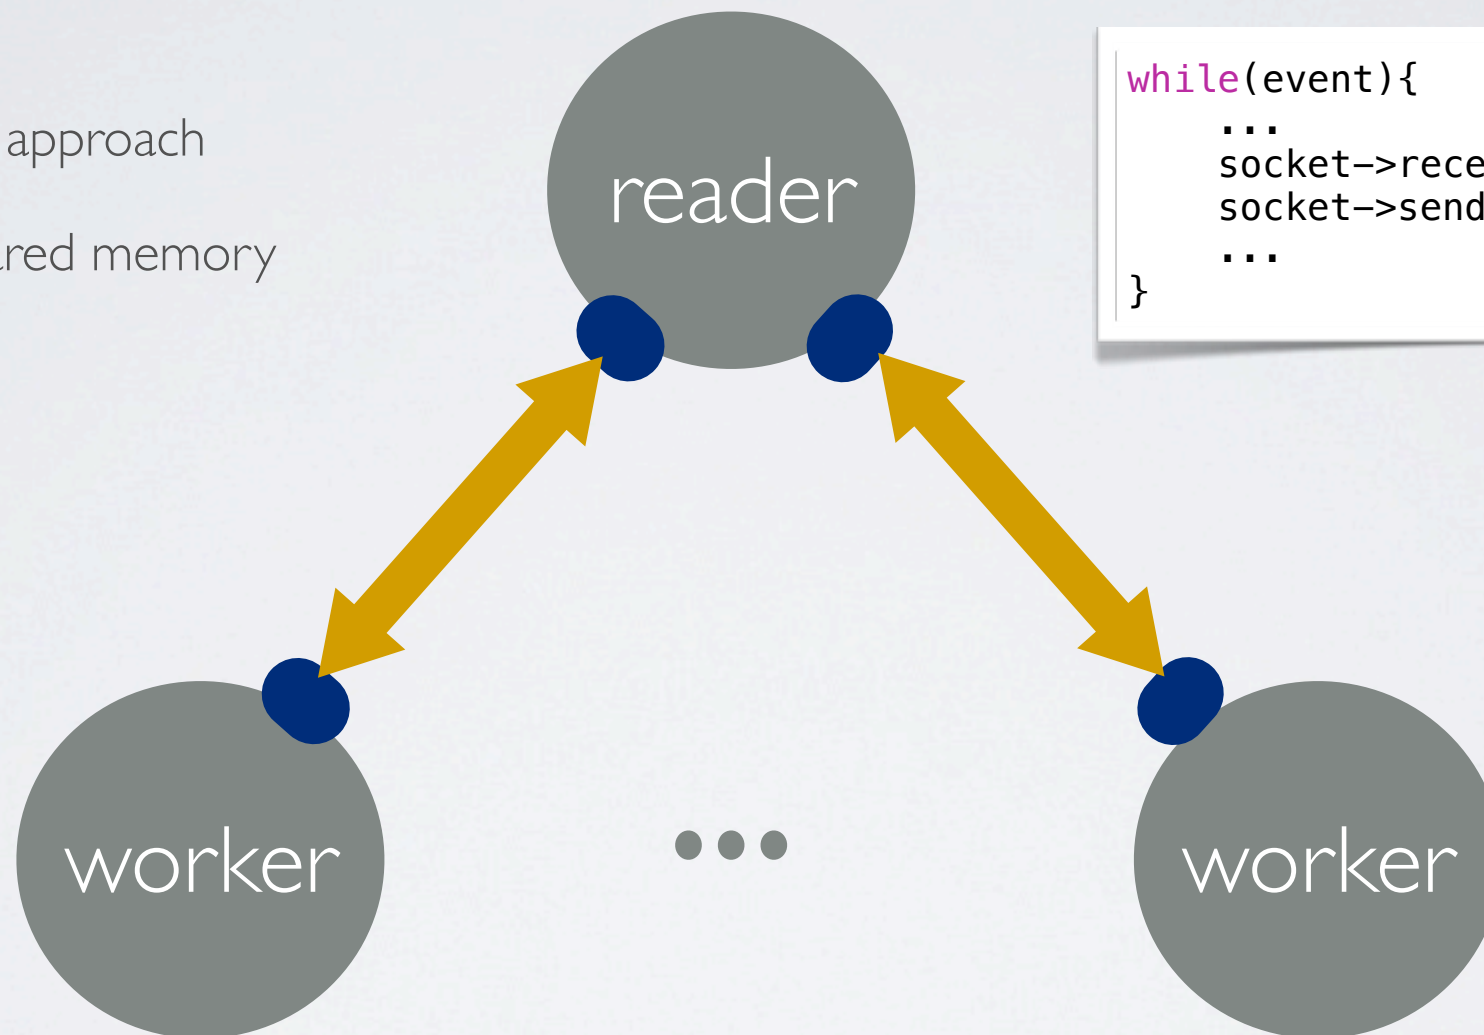
Bandwidth on Linux 3.6



- Benchmark: 1 Megabyte Ping Pong
- Pipe performs 2 copies (U-K-U)
- Spliced pipe performs a single copy (K-U)
- OpenMPI is using shared memory as buffer and performs 2 copies
 - copies are performed in parallel on both ends
 - consumes more CPU
 - pollutes the caches
- ZMQ uses AF_UNIX sockets, i.e. two copies are performed (U-K-U)
 - all sorts of buffers in between
 - optimized for distributed environments

IN ACTION: ATHENAMP

- See Vakho's talk
- Temporary hybrid approach
 - metadata in shared memory
- Benefits:
 - load balancing
 - synchronization



```
while(event){  
    ...  
    socket->receive(&ready_notification);  
    socket->send(event, event_size);  
    ...  
}
```

```
while(true){  
    ...  
    socket->send(ready_notification, size);  
    socket->receive(&event);  
    ...  
}
```

```
while(true){  
    ...  
    socket->send(ready_notification, size);  
    socket->receive(&event);  
    ...  
}
```


CONCLUSION

- Library provides an uniform message-passing abstraction for inter-process communication
- Data and synchronization in a single unit
- Communication patterns and topologies allow to
 - ▶ reduce latency
 - ▶ increase bandwidth
 - ▶ express parallelism
- More implementations will be considered
 - ▶ shared memory segment for small messages (< 10K, low latency)
 - ▶ MPI? Doesn't behave well with forking processes...

QUESTIONS?